

CSCI 2021 Machine Architecture and Organization Fall 2017 Homework Assignment # 3

Instructions:

- This homework must be done individually.
- Posted Tuesday, October 31 and **due on moodle at 11:55 pm on Tuesday, November 14**
- Please **submit as a PDF** containing solutions, moodle will not accept any other format
- There are six problems; we will go over them in discussion section after the due date.
- The textbook in this context is: R. Bryant, D. O'Hallaron, Computer Systems: A Programmer's Perspective. (3rd Edition)

Problem 1 (HCL - similar to Practice Problem 4.11 in the book) (10 points)

Design a logic circuit that finds the second largest value among the set of 3 words, A, B, and C, using an HCL case expression.

```
Word mid3 = [  
  A<=B && A>=C : A;  
  A>=B && A<=C : A;  
  B<=A && B>=C : B;  
  B>=A && B<=C : B;  
  1           : C;  
]
```

Problem 2 (Data Dependencies) (20 points)

The listing below shows a sequence of Y86-64 instructions from a time-critical part of a program. There are five register data dependencies between instructions in this sequence. Fill in the blanks on the right with details about the 5 dependencies: give the number of which instruction (higher numbered) depends on which previous instruction (lower numbered) via which Y86-64 register (that they both access).

| | |
|-----------------------|----------------|
| iaddq \$8,%r8 | #instruction 1 |
| mrmovq 8(%rax), %rcx | #instruction 2 |
| addq %rcx, %rbx | #instruction 3 |
| subq %r8, %rcx | #instruction 4 |
| rrmovq %rcx, %rbx | #instruction 5 |
| rmmovq %rbx, 24(%rax) | #instruction 6 |

- Instruction # 3 depends on instruction # 2 via register %rcx
- Instruction # 4 depends on instruction # 1 via register %r8
- Instruction # 5 depends on instruction # 4 via register %rcx
- Instruction # 6 depends on instruction # 3 via register %rbx
- Instruction # 4 depends on instruction # 2 via register %rcx

f. Only one of these dependencies is a “load-use” hazard that will hurt the performance of the code when running on our pipelined Y86-64 implementation. Which one is that?

Instruction #3

g. You can make the code run faster without changing its behavior by rearranging the instructions so that the “load-use” hazard does not require stalling. Show your improved program after rearranging the instructions.

2 -> 1 -> 3 -> 4 -> 5 -> 6

Problem 3 (Y86 Pipelining) (30 points)

In this exercise, we examine how pipelining affects the clock cycle time of the processor. Assume that individual stages of the data path have the following latencies:

| Fetch | Decode | Execute | Memory | Writeback |
|--------|--------|---------|--------|-----------|
| 250 ps | 500 ps | 150 ps | 250 ps | 200 ps |

Storing the results in an output register costs 20 ps, and each stage costs 20 ps extra for the registers between pipeline stages.

You are working with the following Y86 assembly code:

```
.L1
1    irmovq    $3, %rdx
2    irmovq    $2, %rcx
3    addq      %rcx, %rdx
4    rrmovq    %rcx, %rbx
5    mrmovq    4(%rcx), %rax
6    addq      %rdx, %rax
7    andq      %rax, %rax
8    je        .L1
```

- A. Suppose you are now upgrading yourself to a pipelined processor similar to the one described in Figure 4.41 (p. 424) in the textbook. However, this is a very basic pipelined processor that (1) always predicts branches will not be taken and (2) does not support data forwarding.
- Draw the pipeline stages for each instruction for the first two iterations (i.e., conditional jump is taken) in the following table.** The first two have been done for you. (Hint: You will need to insert “bubbles” to avoid certain hazards.)

Excel spreadsheet attached to the end of document

[illegible]

- B. You have finally decided to treat yourself to a state-of-the-art processor that has 100% branch prediction (don't ask how) and all the data forwarding paths described in Figure 4.2 (p. 440) in the textbook.
- Draw the stages for each instruction for the first two iterations (i.e., conditional jump is taken) in the following table. Also, indicate wherever data forwarding occurs by circling the relevant stages and connecting them with an arrow.**

[illegible]

| | | | | | | | | | | | | | | |
|---|--|--|--|--|--|--|--|--|--|--|--|--|--|--|
| 4 | | | | | | | | | | | | | | |
| 5 | | | | | | | | | | | | | | |
| 6 | | | | | | | | | | | | | | |
| 7 | | | | | | | | | | | | | | |
| 8 | | | | | | | | | | | | | | |
| 1 | | | | | | | | | | | | | | |
| 2 | | | | | | | | | | | | | | |
| 3 | | | | | | | | | | | | | | |
| 4 | | | | | | | | | | | | | | |
| 5 | | | | | | | | | | | | | | |
| 6 | | | | | | | | | | | | | | |
| 7 | | | | | | | | | | | | | | |
| 8 | | | | | | | | | | | | | | |

Problem 4 (Disassembling Code) (15 points)

For each byte sequence listed, determine the Y86 instruction sequence it encodes. There is no invalid byte in any sequence. Your answer should be in the style that is used in the solution of Practice Problem 4.2 (page 360), meaning that you should include the instruction and the memory address of each instruction.

- A. 0x300: 40120f0f0000000000000063221000
- B. 0x400: a05f204530f0feffffffffffff6060b05f90
- C. 0x500: 104065fcffffffffffff10106067741905000000000000000201200

- A.) 0x300 rmmovq %rcx, 0xf0f(%rdx)
0x302 xorq %rdx, %rcx
0x303 nop
0x304 halt
- B.) 0x400 pushq %rbp
0x402 rmmovq %rsp, %rbp
0x40c irmovq \$-2 %rax
0x416 addq %rsi %rax
0x418 popq %rbp
0x41A ret
- C.) 0x500 nop
0x501 rmmovq %rsi -3(%rbp)

```

0x50B nop
0x50C nop
0x50D addq %rsi %rdi
0x50F jne 0x0519
0x510 halt
0x512 rrmovq %rcx %rdx
0x513 halt

```

Problem 5 (x86 to Y86 Conversion) (10 points)

Given the c program:

```

int sum(int x) {
    if (x == 0 || x ==1) {
        return 1;
    } else {
        return x + sum(x-1);
    }
}

```

The following x86-64 assembly code is generated:

```

sum:
    cmpl    $1, %rdi
    ja      .L8
    movl    $1, %eax
    ret

.L8:
    pushq   %rbx
    movl    %edi, %ebx
    leal    -1(%rdi), %edi
    call    sum
    addl    %ebx, %eax
    popq    %rbx
    ret

```

Convert this to Y86-64 assembly code that does the same thing. An example in the textbook can be found on page 365.

```

sum:
    irmovq $1, %r9
    rrmovq %rdi, %r10
    subq %r9, %r10
    jle .L8
    irmovq $1, %rax
    ret

.L8

```

```

pushq %rbx
rrmovq %rdi, %rbx
subq %r9, %rdi
call sum
addq %rbx, %rax
popq %rbx
ret

```

Problem 6 (Y86 Stages) (15 points)

This question describes a new instruction we might consider adding to the Y86-64 processor. Like the textbook does in the section 4.3 and we did in lecture for the existing Y86-64 instructions, give the actions that a SEQ-style Y86-64 implementation would take in each stage to execute the instruction. For each instruction, you should make a table with 6 rows labeled Fetch, Decode, Execute, Memory, Writeback, and PC update (similar to Figure 4.18). Use the same notations as we did in class, with signal names like `valA`, and `R[...]` and `M[...]` to represent accesses to the register file and memory. The `leave` instruction is used to clean up a stack frame when a frame pointer is in use (it's mentioned in section 3.10.5 of the textbook, p. 292). It first copies the frame pointer `%rbp` into the stack pointer `%rsp`, and then it pops the top entry of the stack into `%rbp`. In other words, it is equivalent to the two-instruction sequence:

1. `movq %rbp, %rsp`
2. `popq %rbp`

| | |
|---------|------------------------------------------------------------------------------------------------------------------|
| | Leave |
| Fetch | <code>icode: ifun <- M_1[PC]</code> <code>rA:rB <- M_1[PC + 1]</code> <code>valP <- PC + 2</code> |
| Decode | <code>valA <- R[%rbp]</code> <code>valB <- R[%rsp]</code> |
| Execute | <code>valE <- valB + 8</code> |
| Memory | <code>valM <- m_8[valE]</code> |

| | |
|------------|------------------------------------|
| Write Back | R[%rsp] <- valA R[%rbp] <- valM |
| PC Update | PC <- valP |